

In [133...]

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report
```

In [134...]

```
data=pd.read_csv('Teleco_Customer_Churn.csv')
```

In [135...]

```
data.head()
```

Out[135...]

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Nc
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	...	No	No	No	No	Nc
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	...	Yes	No	No	No	Nc
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	...	No	No	No	No	Nc
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	...	Yes	Yes	No	No	Nc
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	...	No	No	No	No	Nc

5 rows × 21 columns

In [136...]

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   customerID      7043 non-null   object  
 1   gender          7043 non-null   object  
 2   SeniorCitizen   7043 non-null   int64  
 3   Partner         7043 non-null   object  
 4   Dependents     7043 non-null   object  
 5   tenure          7043 non-null   int64  
 6   PhoneService    7043 non-null   object  
 7   MultipleLines   7043 non-null   object  
 8   InternetService 7043 non-null   object  
 9   OnlineSecurity  7043 non-null   object  
 10  OnlineBackup    7043 non-null   object  
 11  DeviceProtection 7043 non-null   object  
 12  TechSupport    7043 non-null   object  
 13  StreamingTV    7043 non-null   object  
 14  StreamingMovies 7043 non-null   object  
 15  Contract        7043 non-null   object  
 16  PaperlessBilling 7043 non-null   object  
 17  PaymentMethod   7043 non-null   object  
 18  MonthlyCharges 7043 non-null   float64 
 19  TotalCharges   7043 non-null   object  
 20  Churn           7043 non-null   object  
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
In [137...]: # to identify no.of rows and columns
data.shape
```

```
Out[137...]: (7043, 21)
```

```
In [138...]: # check for any null vales in data
data.isnull().any()
```

```
Out[138...]: customerID      False
gender          False
SeniorCitizen   False
Partner         False
Dependents     False
tenure          False
PhoneService    False
MultipleLines   False
InternetService False
OnlineSecurity  False
OnlineBackup    False
DeviceProtection False
TechSupport    False
StreamingTV    False
StreamingMovies False
Contract        False
PaperlessBilling False
PaymentMethod   False
MonthlyCharges  False
TotalCharges   False
Churn           False
dtype: bool
```

```
In [139...]: data.isnull().sum()
```

```
Out[139... customerID      0
      gender        0
      SeniorCitizen 0
      Partner        0
      Dependents     0
      tenure         0
      PhoneService    0
      MultipleLines   0
      InternetService 0
      OnlineSecurity   0
      OnlineBackup      0
      DeviceProtection 0
      TechSupport      0
      StreamingTV      0
      StreamingMovies   0
      Contract         0
      PaperlessBilling 0
      PaymentMethod     0
      MonthlyCharges    0
      TotalCharges      0
      Churn            0
      dtype: int64
```

```
In [140... # check for duplicates
duplicates=data.duplicated()
print(duplicates.sum())
```

```
0
```

```
In [141... data.isnull().values.any()
```

```
Out[141... False
```

```
In [142... data['TotalCharges']=pd.to_numeric(data['TotalCharges'],errors='coerce')
total=data['TotalCharges'].sum()
print(total)
```

```
16056168.7
```

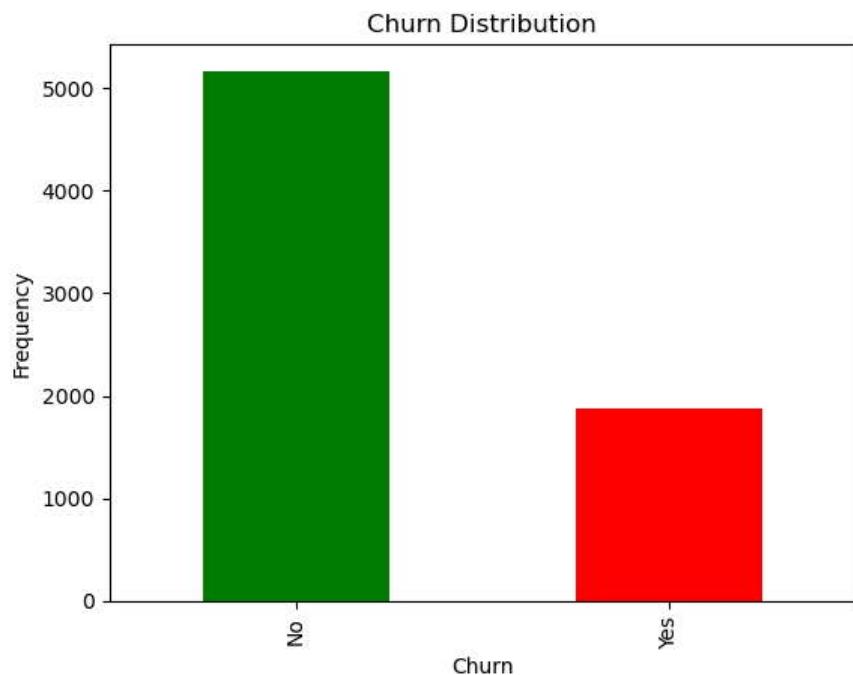
```
In [143... # univariate analysis
# Distribution of churn (Yes/No).
churn_distribution=data['Churn'].value_counts()
print(churn_distribution)
data['Churn'].value_counts().plot(kind='bar', color=['green', 'red'])
plt.title('Churn Distribution')
plt.xlabel('Churn')
plt.ylabel('Frequency')
plt.show()
```

```
Churn
```

```
No      5174
```

```
Yes     1869
```

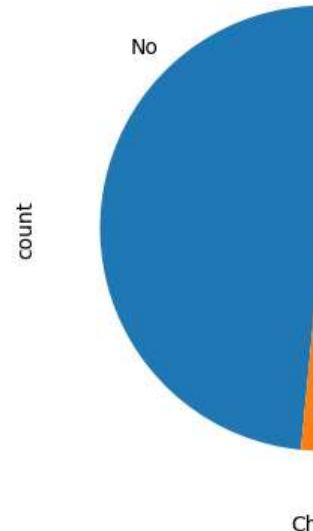
```
Name: count, dtype: int64
```



```
In [144]: # pie chart in terms of percentage
churn_count=data['Churn'].value_counts()
churn_yespercent=(churn_count['Yes']/churn_count.sum())*100
churn_nopercent=(churn_count['No']/churn_count.sum())*100
print(churn_count)
print(f"Churn Yes Percentage: {churn_yespercent:.2f}%")
print(f"Churn No Percentage: {churn_nopercent:.2f}%")
data['Churn'].value_counts().plot(kind='pie',color=['green','red'],labels=churn_count.index)
plt.title('Churn Percentage Distribution')
plt.xlabel('Churn')
plt.show()
```

```
Churn
No      5174
Yes     1869
Name: count, dtype: int64
Churn Yes Percentage: 26.54%
Churn No Percentage: 73.46%
```

## Churn Percentage Distribution



```
In [145...]: #Histograms of numerical features (tenure, MonthlyCharges, TotalCharges).
data[['tenure', 'MonthlyCharges', 'TotalCharges']].info()
plt.figure(figsize=(15, 4))

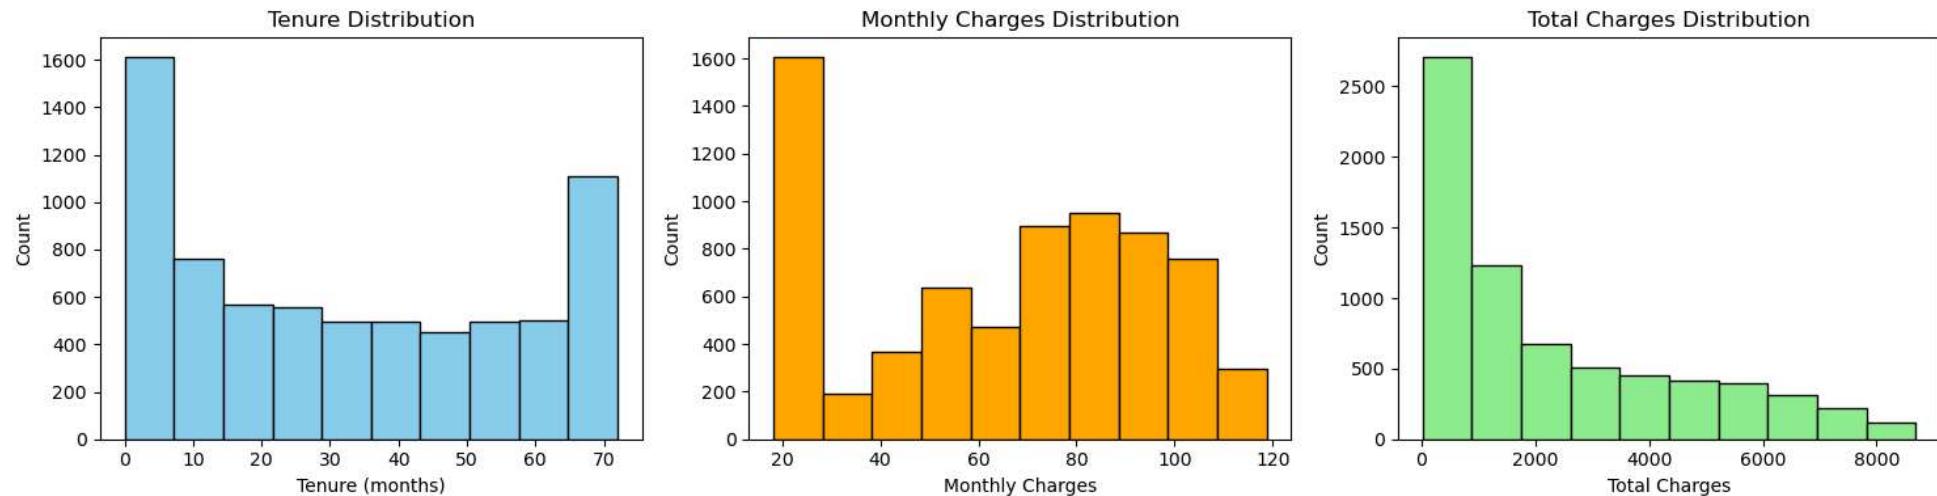
# Tenure
plt.subplot(1, 3, 1)
plt.hist(data['tenure'].dropna(), color='skyblue', edgecolor='black')
plt.title('Tenure Distribution')
plt.xlabel('Tenure (months)')
plt.ylabel('Count')

# Monthly Charges
plt.subplot(1, 3, 2)
plt.hist(data['MonthlyCharges'].dropna(), color='orange', edgecolor='black')
plt.title('Monthly Charges Distribution')
plt.xlabel('Monthly Charges')
plt.ylabel('Count')

# Total Charges
plt.subplot(1, 3, 3)
plt.hist(data['TotalCharges'].dropna(), color='lightgreen', edgecolor='black')
plt.title('Total Charges Distribution')
plt.xlabel('Total Charges')
plt.ylabel('Count')

plt.tight_layout()
plt.show()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 3 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   tenure      7043 non-null   int64  
 1   MonthlyCharges 7043 non-null   float64 
 2   TotalCharges 7032 non-null   float64 
dtypes: float64(2), int64(1)
memory usage: 165.2 KB
```



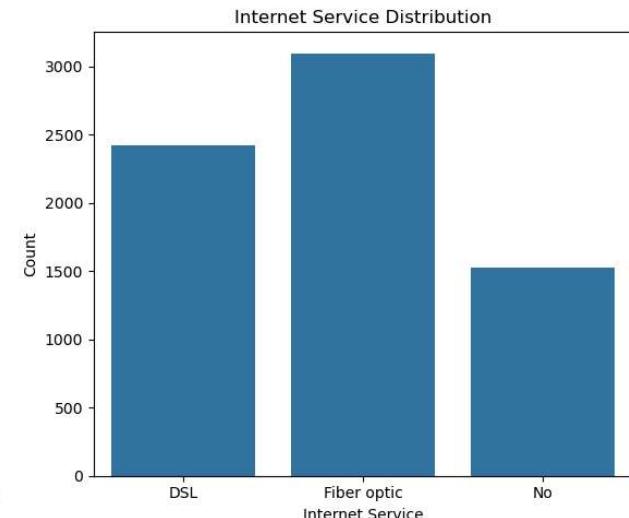
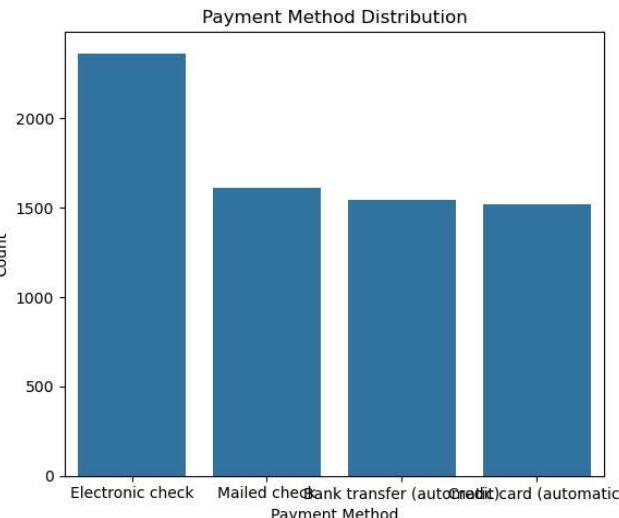
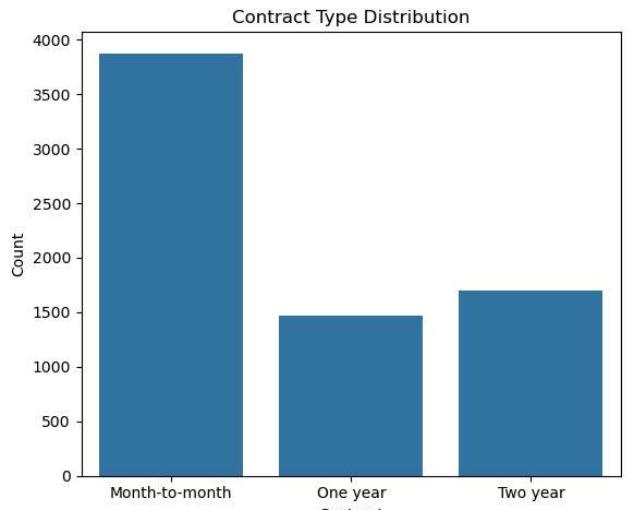
```
In [146...]: # Count plots for categorical features (Contract, PaymentMethod, InternetService).
plt.figure(figsize=(18, 5))
```

```
# Contract
plt.subplot(1, 3, 1)
sns.countplot(data=data, x='Contract')
plt.title('Contract Type Distribution')
plt.xlabel('Contract')
plt.ylabel('Count')

# Payment Method
plt.subplot(1, 3, 2)
sns.countplot(data=data, x='PaymentMethod')
plt.title('Payment Method Distribution')
plt.xlabel('Payment Method')
plt.ylabel('Count')

# Internet Service
plt.subplot(1, 3, 3)
sns.countplot(data=data, x='InternetService')
plt.title('Internet Service Distribution')
plt.xlabel('Internet Service')
plt.ylabel('Count')

plt.tight_layout()
plt.show()
```



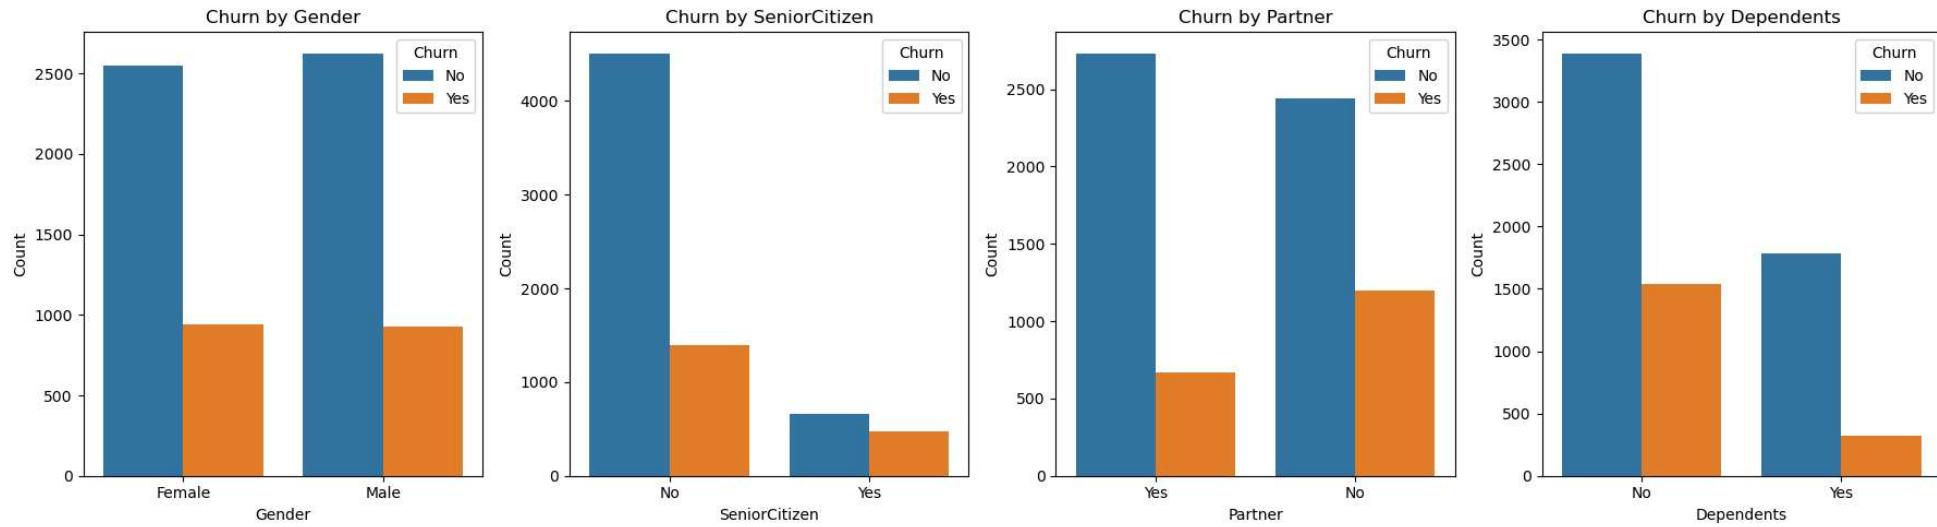
In [147...]

```
# Bivariate Analysis
# Churn by demographics (gender, SeniorCitizen, Partner, Dependents).
plt.figure(figsize=(18, 5))
plt.subplot(1,4,1)
sns.countplot(data=data, x='gender', hue='Churn')
plt.title('Churn by Gender')
plt.xlabel('Gender')
plt.ylabel('Count')

plt.subplot(1,4,2)
data['SeniorCitizen']=data['SeniorCitizen'].map({0:'No',1:'Yes'})
sns.countplot(data=data, x='SeniorCitizen', hue='Churn')
plt.title('Churn by SeniorCitizen')
plt.xlabel('SeniorCitizen')
plt.ylabel('Count')

plt.subplot(1,4,3)
sns.countplot(data=data, x='Partner', hue='Churn')
plt.title('Churn by Partner')
plt.xlabel('Partner')
plt.ylabel('Count')

plt.subplot(1,4,4)
sns.countplot(data=data, x='Dependents', hue='Churn')
plt.title('Churn by Dependents')
plt.xlabel('Dependents')
plt.ylabel('Count')
plt.tight_layout()
plt.show()
```

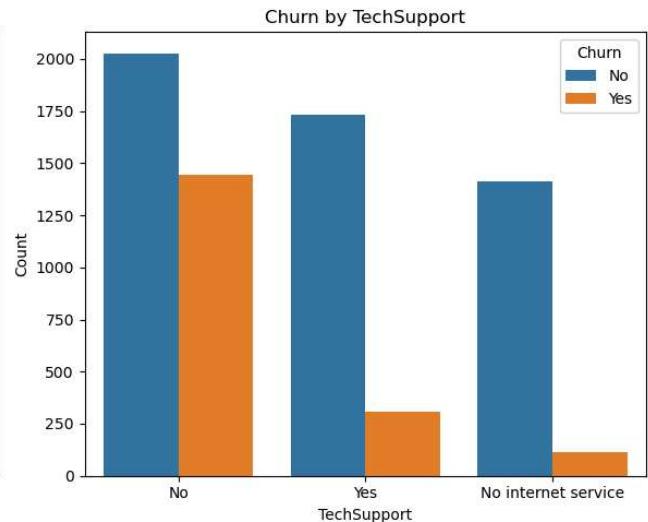
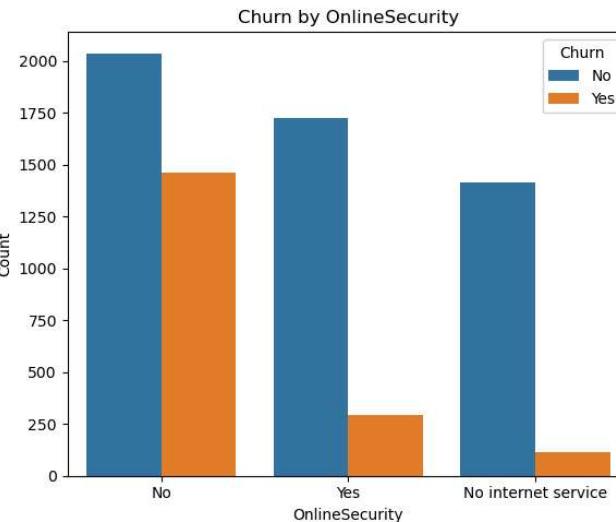
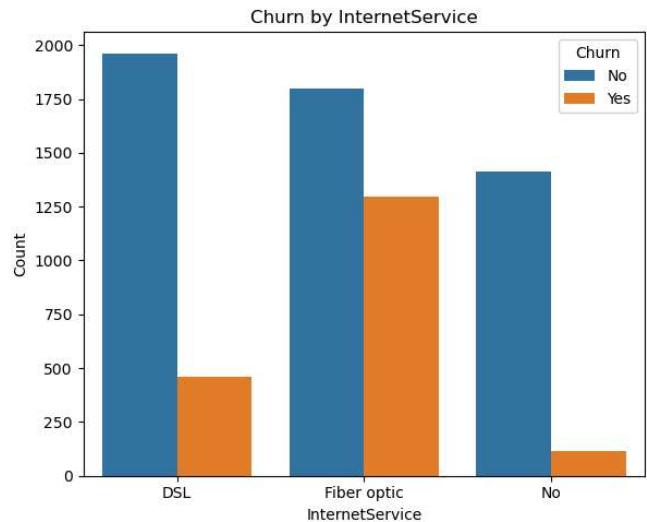


In [148]:

```
#Churn by services (InternetService, OnlineSecurity, TechSupport).
plt.figure(figsize=(18, 5))
plt.subplot(1,3,1)
sns.countplot(data=data, x='InternetService', hue='Churn')
plt.title('Churn by InternetService')
plt.xlabel('InternetService')
plt.ylabel('Count')

plt.subplot(1,3,2)
sns.countplot(data=data, x='OnlineSecurity', hue='Churn')
plt.title('Churn by OnlineSecurity')
plt.xlabel('OnlineSecurity')
plt.ylabel('Count')

plt.subplot(1,3,3)
sns.countplot(data=data, x='TechSupport', hue='Churn')
plt.title('Churn by TechSupport')
plt.xlabel('TechSupport')
plt.ylabel('Count')
plt.tight_layout()
plt.show()
```

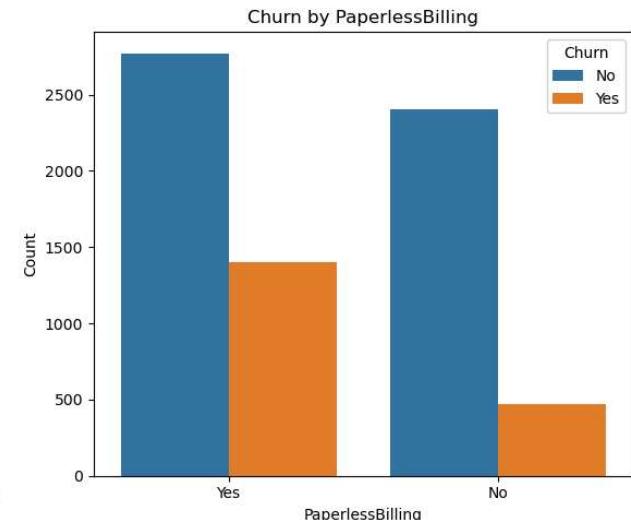
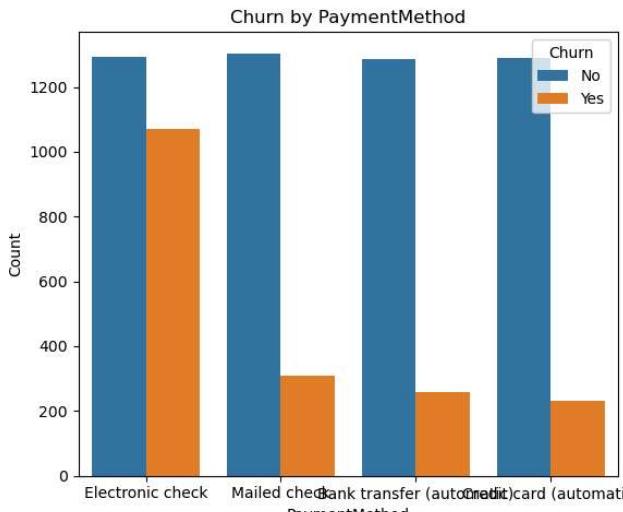
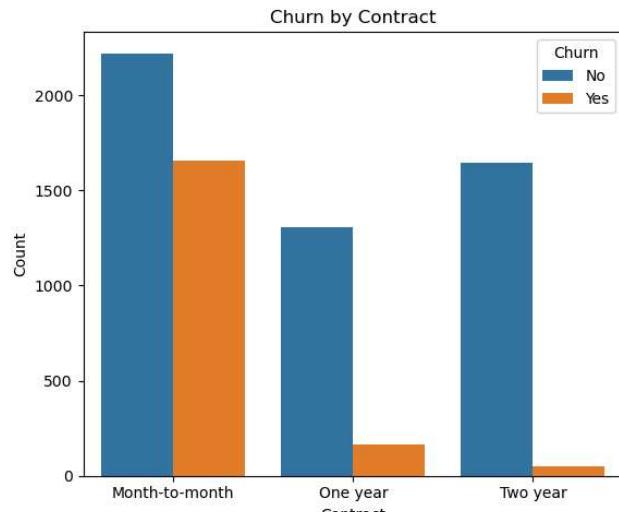


In [149]:

```
#Churn by account details (Contract, PaymentMethod, PaperlessBilling).
plt.figure(figsize=(18, 5))
plt.subplot(1,3,1)
sns.countplot(data=data, x='Contract', hue='Churn')
plt.title('Churn by Contract')
plt.xlabel('Contract')
plt.ylabel('Count')

plt.subplot(1,3,2)
sns.countplot(data=data, x='PaymentMethod', hue='Churn')
plt.title('Churn by PaymentMethod')
plt.xlabel('PaymentMethod')
plt.ylabel('Count')

plt.subplot(1,3,3)
sns.countplot(data=data, x='PaperlessBilling', hue='Churn')
plt.title('Churn by PaperlessBilling')
plt.xlabel('PaperlessBilling')
plt.ylabel('Count')
plt.tight_layout()
plt.show()
```

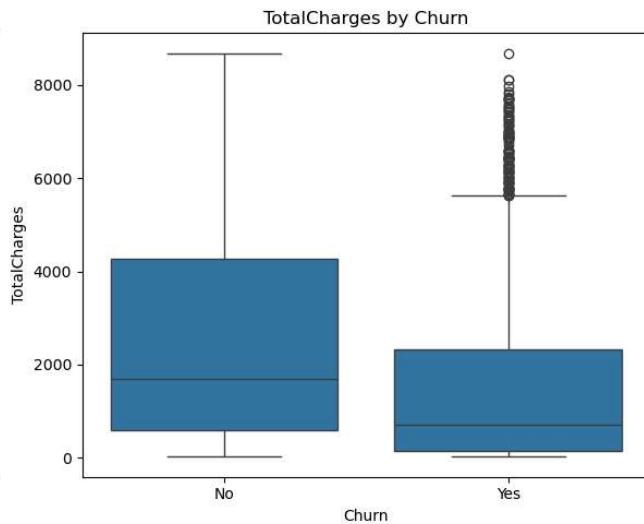
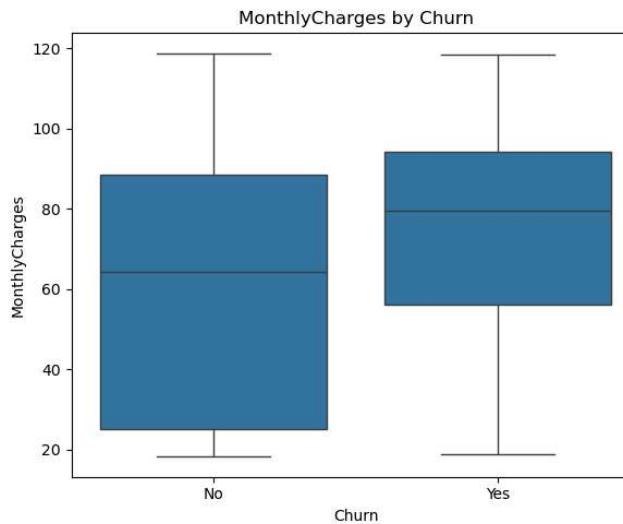
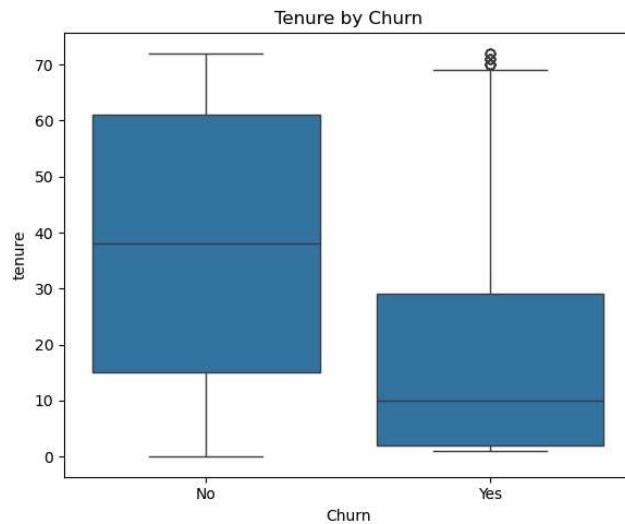


In [150...]

```
#Box plots of tenure & charges by churn
plt.figure(figsize=(18, 5))
plt.subplot(1,3,1)
sns.boxplot(data=data, x='Churn', y='tenure')
plt.title('Tenure by Churn')

plt.subplot(1,3,2)
sns.boxplot(data=data, x='Churn', y='MonthlyCharges')
plt.title('MonthlyCharges by Churn')

plt.subplot(1,3,3)
sns.boxplot(data=data, x='Churn', y='TotalCharges')
plt.title('TotalCharges by Churn')
plt.tight_layout()
plt.show()
```



In [151...]

```
#KDE(Kernel Density Estimate) plots of tenure & charges by churn
```

```
plt.figure(figsize=(18,5))
plt.subplot(1,3,1)
```

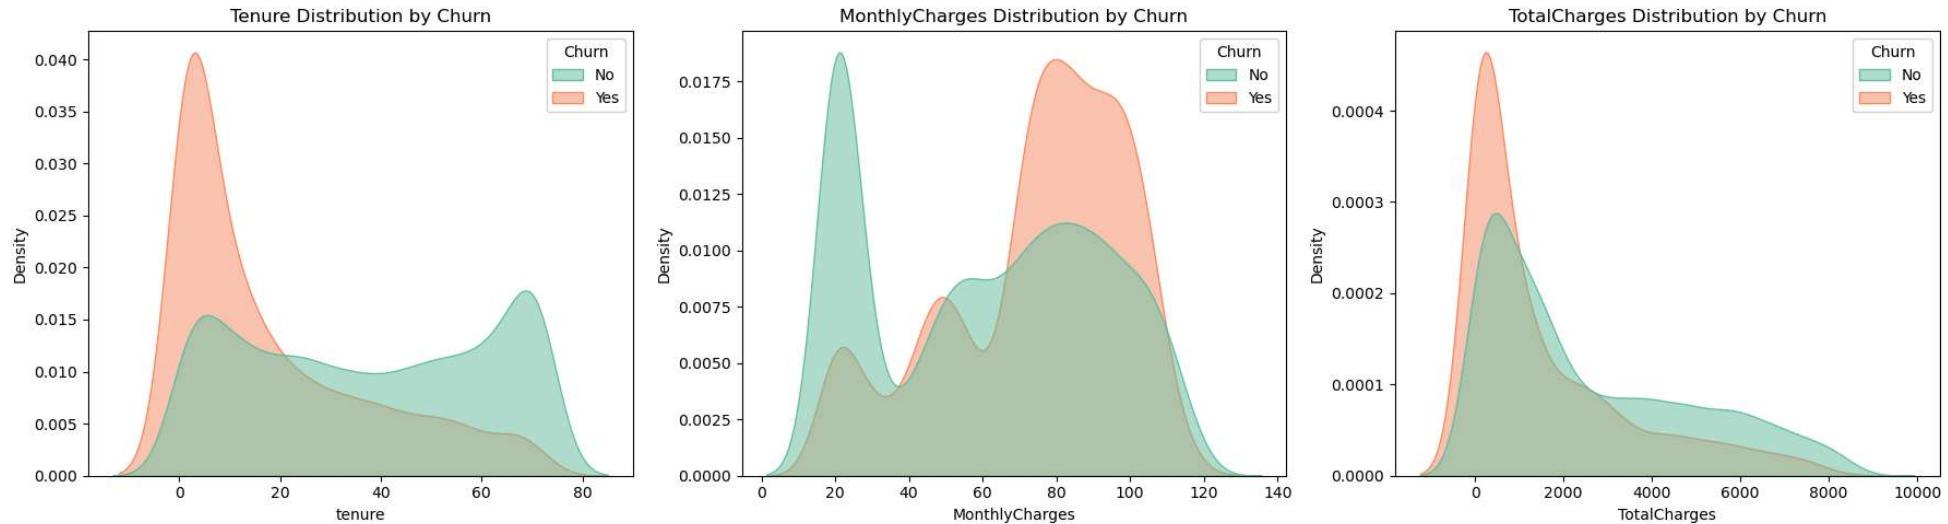
```

sns.kdeplot(data=data, x='tenure', hue='Churn', fill=True, common_norm=False, palette='Set2', alpha=0.5)
plt.title('Tenure Distribution by Churn')

plt.subplot(1,3,2)
sns.kdeplot(data=data, x='MonthlyCharges', hue='Churn', fill=True, common_norm=False, palette='Set2', alpha=0.5)
plt.title('MonthlyCharges Distribution by Churn')

plt.subplot(1,3,3)
sns.kdeplot(data=data, x='TotalCharges', hue='Churn', fill=True, common_norm=False, palette='Set2', alpha=0.5)
plt.title('TotalCharges Distribution by Churn')
plt.tight_layout()
plt.show()

```



In [152]: #Scatter plot: tenure vs. TotalCharges (color by churn).

```

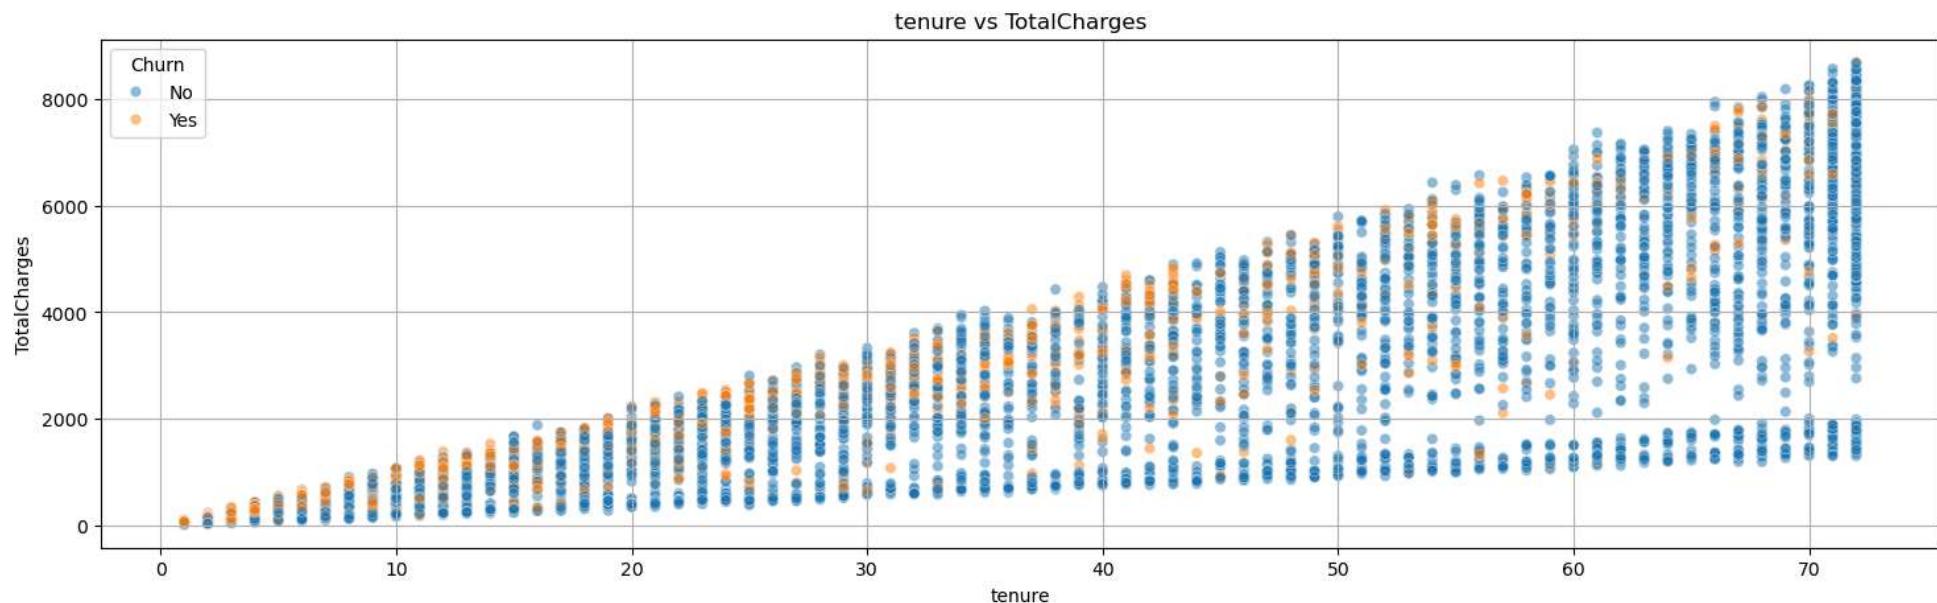
plt.figure(figsize=(18,5))
sns.scatterplot(data=data,x='tenure',y='TotalCharges',hue='Churn',alpha=0.5)
plt.title('tenure vs TotalCharges')
plt.xlabel('tenure')
plt.ylabel('TotalCharges')
plt.grid(True)
plt.legend(title='Churn')
plt.show()

```

#Upward trend: Generally, longer tenure → higher total charges.

#Churn clusters: churned customers clustered in short-tenure, low-charge regions.

#Outliers: Look for any odd points (e.g., long-tenure with low charges).



In [153...]

```
#Multivariate & Advanced Analysis
#Correlation heatmap of numerical features
data['Churn'] = data['Churn'].map({'Yes': 1, 'No': 0})
numeric_columns = data.select_dtypes(include=['float64', 'int64'])
print(numeric_columns)
corr_matrix = numeric_columns.corr()
print(corr_matrix)
plt.figure(figsize=(10, 6))
sns.heatmap(
    corr_matrix,
    annot=True,           # Show correlation values
    fmt='.2f',            # Format numbers to 2 decimal places
    cmap='coolwarm',      # Color map (blue = negative, red = positive)
    square=True,
    linewidths=0.5,
    cbar_kws={'shrink': 0.7} # Shrink colorbar
)
plt.title('Correlation Heatmap of Numerical Features')
plt.tight_layout()
plt.show()

#Values near 1 → strong positive correlation
#Values near -1 → strong negative correlation
#Values near 0 → little or no correlation
#tenure and TotalCharges Likely have a strong positive correlation
#MonthlyCharges and TotalCharges might also be somewhat correlated
#weak or no correlation with variables like SeniorCitizen
```

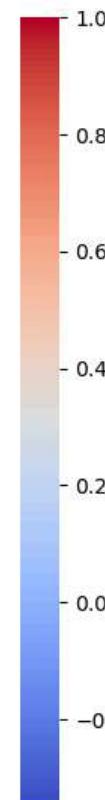
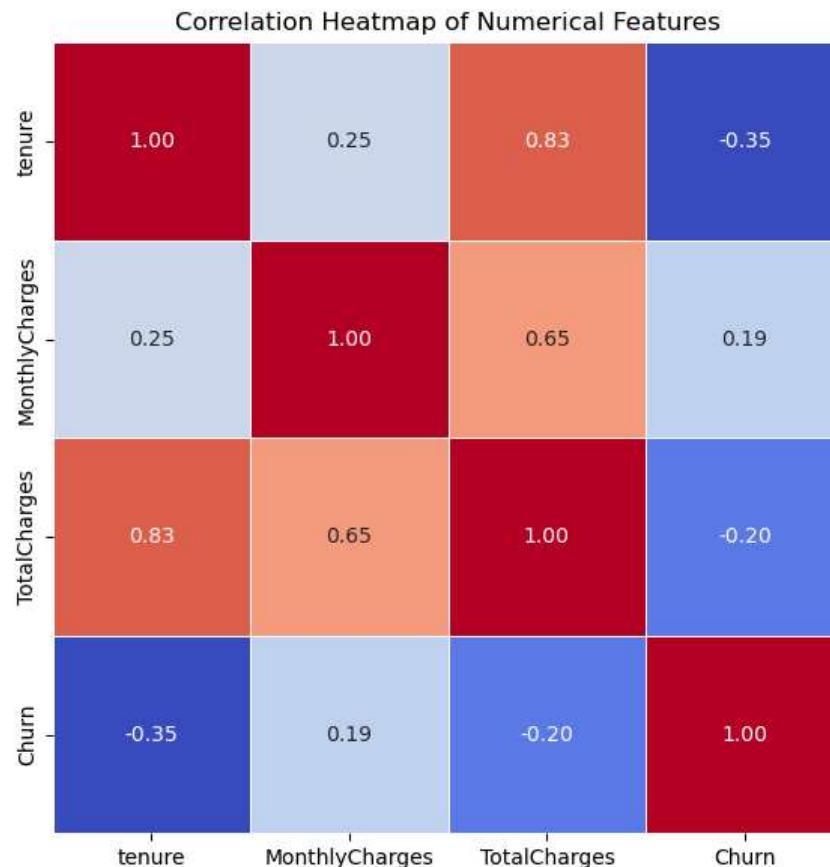
```

tenure  MonthlyCharges  TotalCharges  Churn
0        1              29.85        29.85    0
1        34             56.95        1889.50   0
2        2              53.85        108.15    1
3        45             42.30        1840.75   0
4        2              70.70        151.65    1
...
7038    24             84.80        1990.50   0
7039    72             103.20        7362.90   0
7040    11             29.60        346.45    0
7041    4              74.40        306.60    1
7042    66             105.65        6844.50   0

```

[7043 rows x 4 columns]

	tenure	MonthlyCharges	TotalCharges	Churn
tenure	1.000000	0.247900	0.825880	-0.352229
MonthlyCharges	0.247900	1.000000	0.651065	0.193356
TotalCharges	0.825880	0.651065	1.000000	-0.199484
Churn	-0.352229	0.193356	-0.199484	1.000000



```
In [154]: #Interaction effects: e.g., MonthlyCharges x InternetService type on churn
plt.figure(figsize=(10, 6))
```

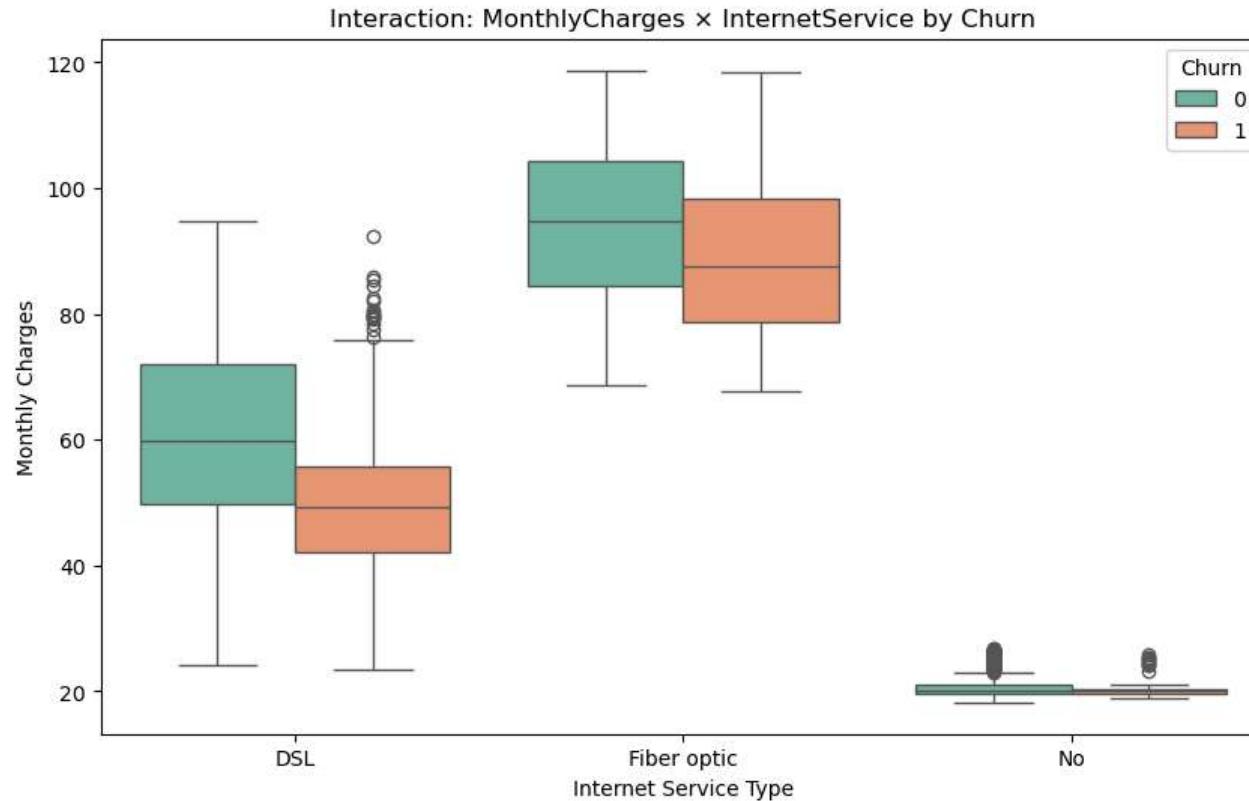
```

sns.boxplot(
    data=data,
    x='InternetService',
    y='MonthlyCharges',
    hue='Churn',
    palette='Set2'
)

```

```
)  
plt.title('Interaction: MonthlyCharges x InternetService by Churn')  
plt.ylabel('Monthly Charges')  
plt.xlabel('Internet Service Type')  
plt.legend(title='Churn')  
plt.show()
```

#Compare how MonthlyCharges varies for churned vs. not churned customers within each Internet type (DSL, Fiber optic, No).  
#higher churn with Fiber optic + high charges.



In [ ]:

```
In [155...]:  
#High-risk churn customer segments profile.  
#data['Churn'] = data['Churn'].map({'Yes': 1, 'No': 0})  
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   customerID      7043 non-null   object  
 1   gender          7043 non-null   object  
 2   SeniorCitizen   7043 non-null   object  
 3   Partner         7043 non-null   object  
 4   Dependents     7043 non-null   object  
 5   tenure          7043 non-null   int64  
 6   PhoneService    7043 non-null   object  
 7   MultipleLines   7043 non-null   object  
 8   InternetService 7043 non-null   object  
 9   OnlineSecurity  7043 non-null   object  
 10  OnlineBackup    7043 non-null   object  
 11  DeviceProtection 7043 non-null   object  
 12  TechSupport    7043 non-null   object  
 13  StreamingTV    7043 non-null   object  
 14  StreamingMovies 7043 non-null   object  
 15  Contract        7043 non-null   object  
 16  PaperlessBilling 7043 non-null   object  
 17  PaymentMethod   7043 non-null   object  
 18  MonthlyCharges 7043 non-null   float64 
 19  TotalCharges   7032 non-null   float64 
 20  Churn          7043 non-null   int64  
dtypes: float64(2), int64(2), object(17)
memory usage: 1.1+ MB
```

```
In [156...]: data['Churn'] = pd.to_numeric(data['Churn'], errors='coerce')
print(data['Churn'].isna().sum())
```

```
0
```

```
In [157...]: churn_rate = data['Churn'].mean()
print(f"Overall Churn Rate: {churn_rate*100:.2f}%")
#~26-27% churn in this dataset
print("\nChurn by Contract:")
print(data.groupby('Contract')['Churn'].mean().sort_values(ascending=False))
#Month-to-month customers usually have the highest churn (~43%),
#while Two year contracts have lowest churn (~2-3%).
#Fiber optic customers churn much more than DSL users - Likely due to pricing or service quality.
print("\nChurn by Internet Service:")
print(data.groupby('InternetService')['Churn'].mean().sort_values(ascending=False))
#Customers with <1 year tenure churn the most - new customers are more at risk.
bins = [0, 12, 24, 48, 72]
labels = ['<1 year', '1-2 years', '2-4 years', '4-6 years']
data['tenure_group'] = pd.cut(data['tenure'], bins=bins, labels=labels, include_lowest=True)
print(data.groupby('tenure_group')['Churn'].mean().sort_values(ascending=False))
#Customers paying by Electronic check are often the most likely to churn.
print(data.groupby('PaymentMethod')['Churn'].mean().sort_values(ascending=False))
#High monthly charges → higher churn risk.
data['MonthlyChargesRange'] = pd.cut(data['MonthlyCharges'], bins=[0, 35, 70, 120], labels=['Low', 'Medium', 'High'])
print(data.groupby('MonthlyChargesRange')['Churn'].mean().sort_values(ascending=False))
```

Overall Churn Rate: 26.54%

Churn by Contract:

Contract  
Month-to-month 0.427097  
One year 0.112695  
Two year 0.028319  
Name: Churn, dtype: float64

Churn by Internet Service:

InternetService  
Fiber optic 0.418928  
DSL 0.189591  
No 0.074050  
Name: Churn, dtype: float64

tenure\_group

<1 year 0.474382  
1-2 years 0.287109  
2-4 years 0.203890  
4-6 years 0.095132

Name: Churn, dtype: float64

PaymentMethod

Electronic check 0.452854  
Mailed check 0.191067  
Bank transfer (automatic) 0.167098  
Credit card (automatic) 0.152431

Name: Churn, dtype: float64

MonthlyChargesRange

High 0.353614  
Medium 0.239420  
Low 0.108934

Name: Churn, dtype: float64

C:\Users\kalia\AppData\Local\Temp\ipykernel\_4204\13899394.py:17: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
print(data.groupby('tenure_group')['Churn'].mean().sort_values(ascending=False))
```

C:\Users\kalia\AppData\Local\Temp\ipykernel\_4204\13899394.py:23: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
print(data.groupby('MonthlyChargesRange')['Churn'].mean().sort_values(ascending=False))
```

In [158...]

# / Segment	/ Typical Behavior	/ Churn Rate
# / -----	/ -----	/ -----
# / **Contract: Month-to-month**	/ Short-term, less commitment	/ Very High
# / **Internet: Fiber optic**	/ High-speed users, often price-sensitive	/ High
# / **Payment: Electronic check**	/ Possibly less engaged, prefer quick cancellation	/ High
# / **Tenure: <1 year**	/ New customers, low loyalty	/ Highest
# / **High MonthlyCharges**	/ Expensive plans, potential dissatisfaction	/ High
# / **Senior Citizens (optional)**	/ May find services complex or costly	/ Moderate to High
# High-risk churn customers are typically:		
# On month-to-month contracts		
# Using fiber optic internet		
# Paying via electronic check		
# Having tenure under 1 year		
# Paying high monthly charges		
# Often without online security or tech support add-ons		

In [159...]

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 23 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   customerID        7043 non-null   object  
 1   gender             7043 non-null   object  
 2   SeniorCitizen     7043 non-null   object  
 3   Partner            7043 non-null   object  
 4   Dependents         7043 non-null   object  
 5   tenure             7043 non-null   int64  
 6   PhoneService       7043 non-null   object  
 7   MultipleLines      7043 non-null   object  
 8   InternetService   7043 non-null   object  
 9   OnlineSecurity     7043 non-null   object  
 10  OnlineBackup       7043 non-null   object  
 11  DeviceProtection  7043 non-null   object  
 12  TechSupport        7043 non-null   object  
 13  StreamingTV        7043 non-null   object  
 14  StreamingMovies    7043 non-null   object  
 15  Contract           7043 non-null   object  
 16  PaperlessBilling  7043 non-null   object  
 17  PaymentMethod      7043 non-null   object  
 18  MonthlyCharges    7043 non-null   float64 
 19  TotalCharges       7032 non-null   float64 
 20  Churn              7043 non-null   int64  
 21  tenure_group       7043 non-null   category 
 22  MonthlyChargesRange 7043 non-null   category 
dtypes: category(2), float64(2), int64(2), object(17)
memory usage: 1.1+ MB
```

In [160...]

```
# Load the raw dataset
data = pd.read_csv("Teleco_Customer_Churn.csv")

# Inspect the Churn column
print("Unique values before cleaning:", data['Churn'].unique())
# Keep only valid churn entries ('Yes'/'No')
data = data[data['Churn'].isin(['Yes', 'No'])]

# Encode target variable
data['Churn'] = data['Churn'].map({'Yes': 1, 'No': 0})

# Encode gender
data['gender'] = LabelEncoder().fit_transform(data['gender'])

# Drop customerID
X = data.drop(columns=['customerID', 'Churn'])
y = data['Churn']

# One-hot encode categorical columns
X = pd.get_dummies(X, drop_first=True)

# Split into train/test
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

print("\nDataset split completed successfully!")
print(f"Training samples: {X_train.shape[0]}")
print(f"Testing samples: {X_test.shape[0]}")
print(f"Churn ratio: {y.mean():.2%}")
```

Unique values before cleaning: ['No' 'Yes']

Dataset split completed successfully!

Training samples: 5634

Testing samples: 1409

Churn ratio: 26.54%

In [161...]

```
# Train Random Forest
rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)

# Make predictions
y_pred = rf.predict(X_test)

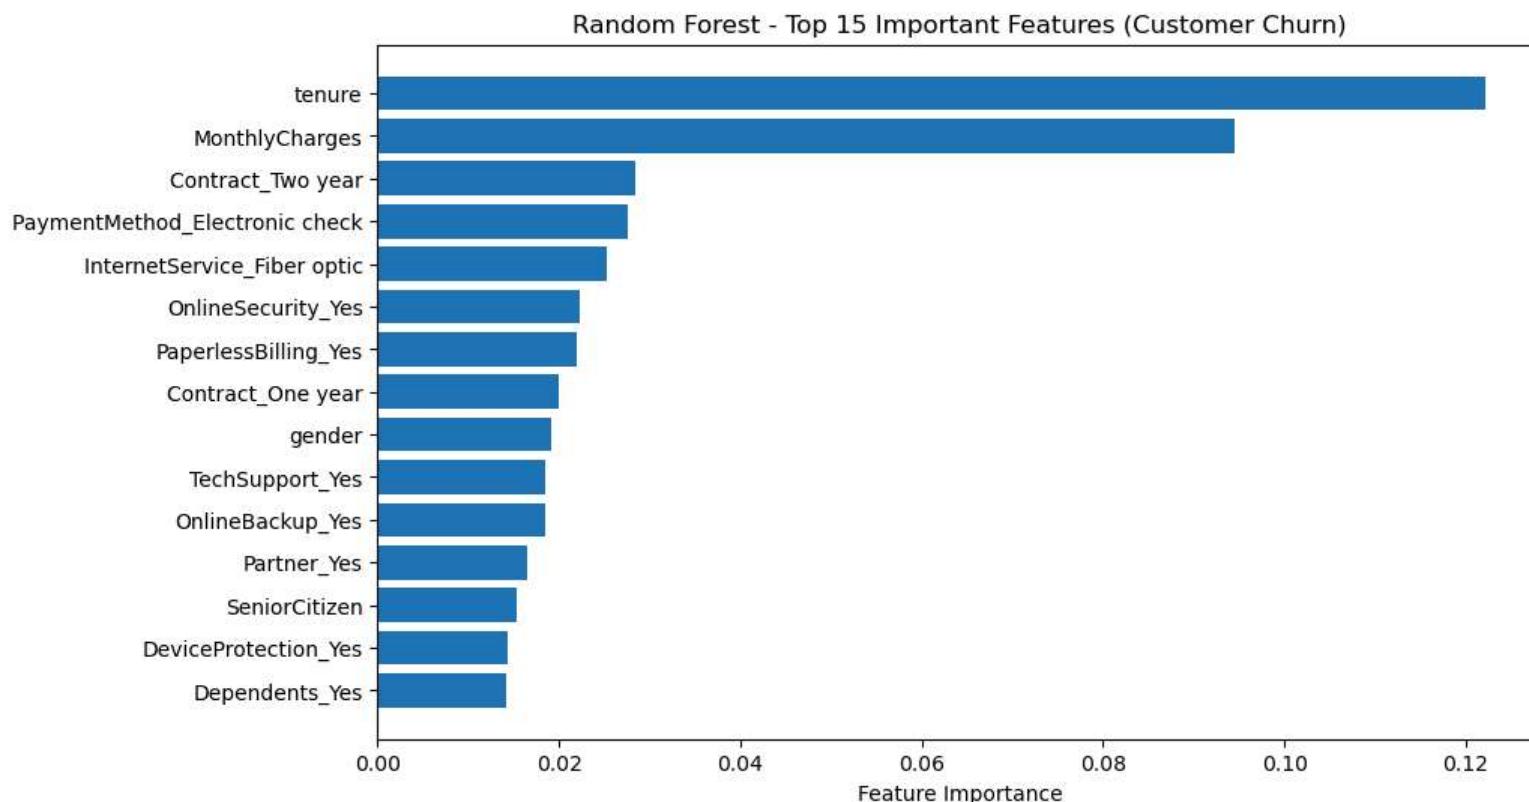
# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)

print(f"Random Forest Accuracy: {accuracy * 100:.2f}%")

# Get feature importance
rf_importance = pd.DataFrame({
    'Feature': X.columns,
    'Importance': rf.feature_importances_
}).sort_values(by='Importance', ascending=False)

# Plot
plt.figure(figsize=(10, 6))
plt.barh(rf_importance['Feature'][:15], rf_importance['Importance'][:15])
plt.xlabel('Feature Importance')
plt.title('Random Forest - Top 15 Important Features (Customer Churn)')
plt.gca().invert_yaxis()
plt.show()
```

Random Forest Accuracy: 79.99%



In [162]

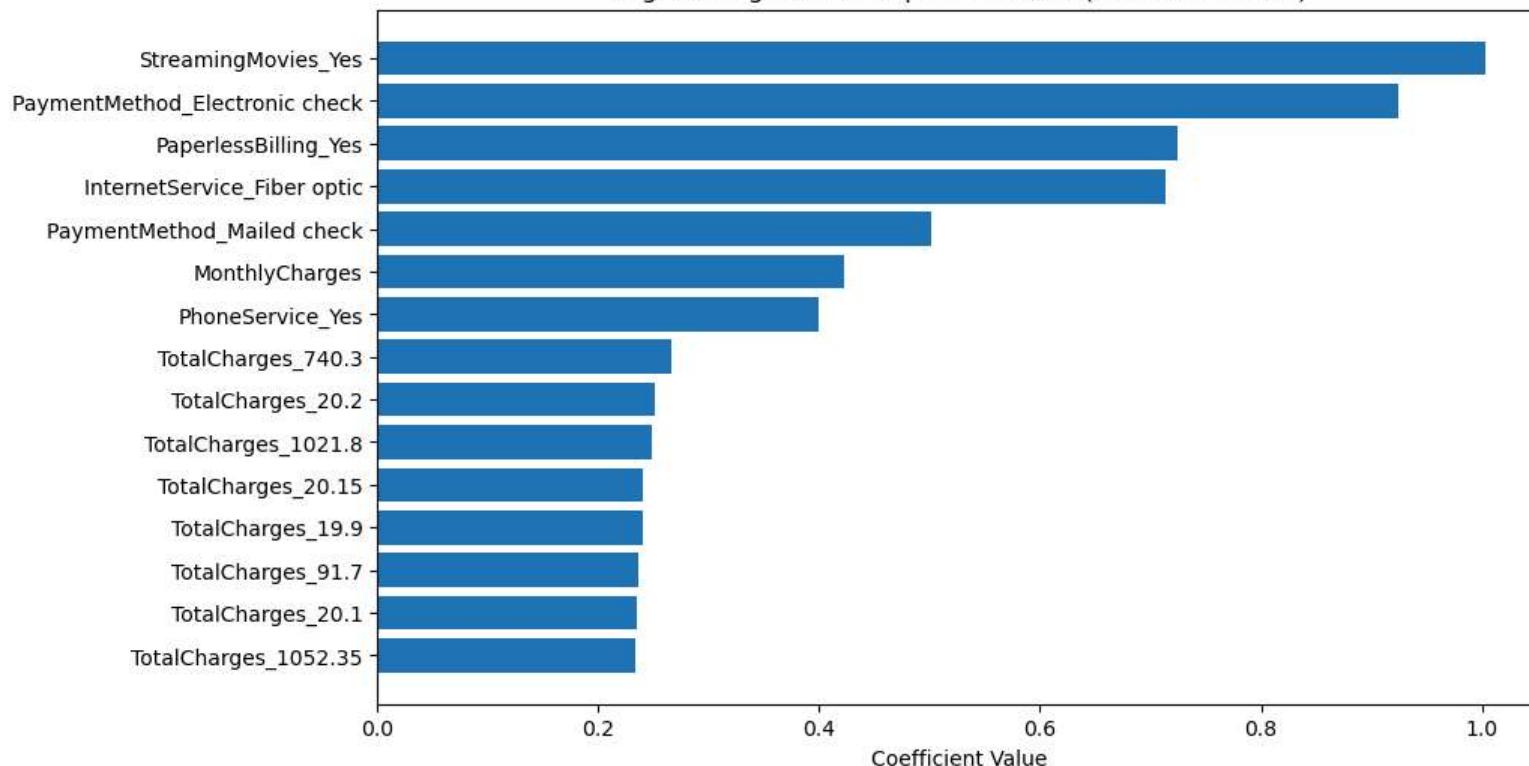
```
# Scale features (important for Logistic Regression)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Train Logistic Regression
lr = LogisticRegression(max_iter=1000)
lr.fit(X_train_scaled, y_train)

# Get feature coefficients
lr_importance = pd.DataFrame({
    'Feature': X.columns,
    'Coefficient': lr.coef_[0]
}).sort_values(by='Coefficient', ascending=False)

# Plot
plt.figure(figsize=(10, 6))
plt.barh(lr_importance['Feature'][:15], lr_importance['Coefficient'][:15])
plt.xlabel('Coefficient Value')
plt.title('Logistic Regression - Top 15 Features (Customer Churn)')
plt.gca().invert_yaxis()
plt.show()
```

### Logistic Regression - Top 15 Features (Customer Churn)



In [163]:

```
top_rf = rf_importance.head(10)
top_lr = lr_importance.head(10)

print("Top Random Forest Features:\n", top_rf)
print("Top Logistic Regression Features:\n", top_lr)
```

Top Random Forest Features:

	Feature	Importance
2	tenure	0.122216
3	MonthlyCharges	0.094510
24	Contract_Two year	0.028415
27	PaymentMethod_Electronic check	0.027663
9	InternetService_Fiber optic	0.025230
12	OnlineSecurity_Yes	0.022367
25	PaperlessBilling_Yes	0.022026
23	Contract_One year	0.019967
0	gender	0.019181
18	TechSupport_Yes	0.018548

Top Logistic Regression Features:

	Feature	Coefficient
22	StreamingMovies_Yes	1.003931
27	PaymentMethod_Electronic check	0.925098
25	PaperlessBilling_Yes	0.724422
9	InternetService_Fiber optic	0.713462
28	PaymentMethod_Mailed check	0.500785
3	MonthlyCharges	0.423186
6	PhoneService_Yes	0.399536
5755	TotalCharges_740.3	0.266454
1638	TotalCharges_20.2	0.250553
71	TotalCharges_1021.8	0.247891

```
In [164...]
#Predict churn using Logistic Regression; evaluate with confusion matrix
#Train Logistic Regression model
log_reg = LogisticRegression(max_iter=1000, random_state=42)
log_reg.fit(X_train, y_train)

#Make predictions
y_pred = log_reg.predict(X_test)

#Evaluate performance
print(" Model Evaluation:")
print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

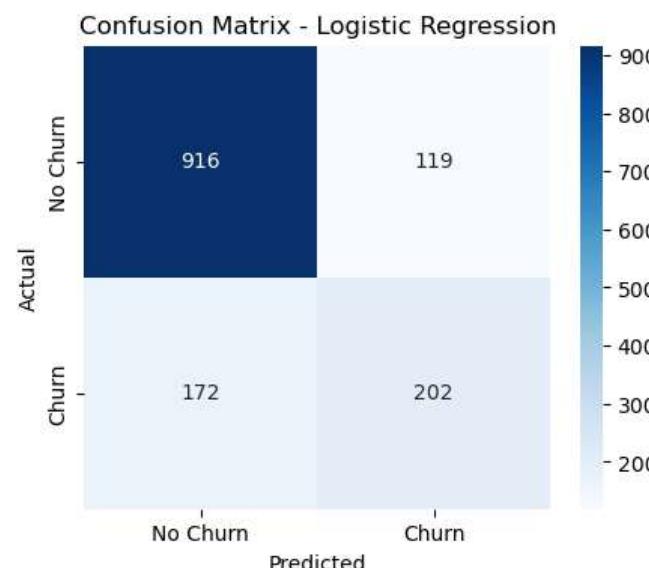
# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['No Churn', 'Churn'], yticklabels=['No Churn', 'Churn'])
plt.title('Confusion Matrix - Logistic Regression')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

Model Evaluation:

Accuracy: 0.7935

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.89	0.86	1035
1	0.63	0.54	0.58	374
accuracy			0.79	1409
macro avg	0.74	0.71	0.72	1409
weighted avg	0.79	0.79	0.79	1409



In [165...]

# Key Findings:

# Top predictors of churn:

```
# Tenure (shorter = higher churn risk)

# MonthlyCharges (higher = higher churn risk)

# Contract type (month-to-month contracts show higher churn)

# Payment method (electronic check customers churn more)

# Lack of Online Security or Tech Support services increases churn risk
```